

Rebuilding optimizing compiler for Dart

{ @mraleph | vegorov@google.com }

6 years of Dart (VM engineer perspective)

{ @mraleph | vegorov@google.com }

Excelsior JET

v8

Dart VM

LuaJIT



Dart VM (2012 - ...)

Dart 2011

«Dart targets a wide range of development scenarios: from a one-person project without much structure to a large-scale project needing formal types in the code to state programmer intent. *To support this wide range of projects, Dart has optional types; this means you can start coding without types and add them later as needed.* We believe Dart will be great for writing large web applications.»

— Dart: A language for structured web programming

```
Dog d = new Cat();  
d.woof();
```

```
Dog d = new Cat();
d.woof();

class Cat {
    noSuchMethod(invocation) {
        if (invocation.memberName == '#woof') {
            print('I am a 🐱 not a 🐶!');
        }
    }
}
```

dynamic language

UUTT

initially
AST based

2012

started new JIT compiler

inline caching (IC)

+

speculative opts

```
// call-site specific cache
ICData { class ↞ (target, frequency) }

dog.woof();
// generic dispatch stub
movq rcx,[rbx+0x27]
inc dword ptr [rcx+0x87]
mov rax, dword ptr [thr+0x20]
cmp byte ptr [rax+0x60], 0
jnz 0x0000000108a41c8c
mov r10, dword ptr [rbx+0x17]
mov r13, dword ptr [rbx+0x7]
```

```
// call-site specific cache
ICData { class ↞ (target, frequency) }

dog.woof();
// generic dispatch stub
dispatch(icData, receiver, ...) {
    target = icData.lookup(receiver.classId)
    if (target == null) {
        target = LookupAndCache(icData, receiver)
    }
    invoke target(...)
}
```

```
    └─ {}  
dog.woof();
```

```
    └─ {}  
dog.woof();  
└─  
  class Dog {  
    woof() => print("🐶🐶🐶");  
  }
```

```
    {}  
    ↪LookupAndCache()  
dog.woof();  
}  
  
class Dog {  
    woof() => print("🐶🐶🐶");  
}
```

```
{}  
↳LookupAndCache()  
dog.woof();  
  
class Dog {  
    woof() => print("🐶🐶🐶");  
}
```

```
    └─ {Dog ↪ (Dog.woof, 1)}  
dog.woof();  
{  
  class Dog {  
    woof() => print("🐶🐶🐶");  
  }  
}
```

```
    └─ {Dog ↪ (Dog.woof, 2)}  
dog.woof();  
{  
  class Dog {  
    woof() => print("🐶🐶🐶");  
  }  
}
```

```
    └─ {Dog ↪ (Dog.woof, 3)}  
dog.woof();  
{  
  class Dog {  
    woof() => print("🐶🐶🐶");  
  }  
}
```

```
    └─ {Dog ↪ (Dog.woof, 4)}  
dog.woof();  
{  
  class Dog {  
    woof() => print("🐶🐶🐶");  
  }  
}
```

```
    └─ {Dog ↪ (Dog.woof, 4)}  
dog.woof();  
{  
  └─ class Wolf {  
      woof() => print("🐺🐺🐺");  
    }  
}
```

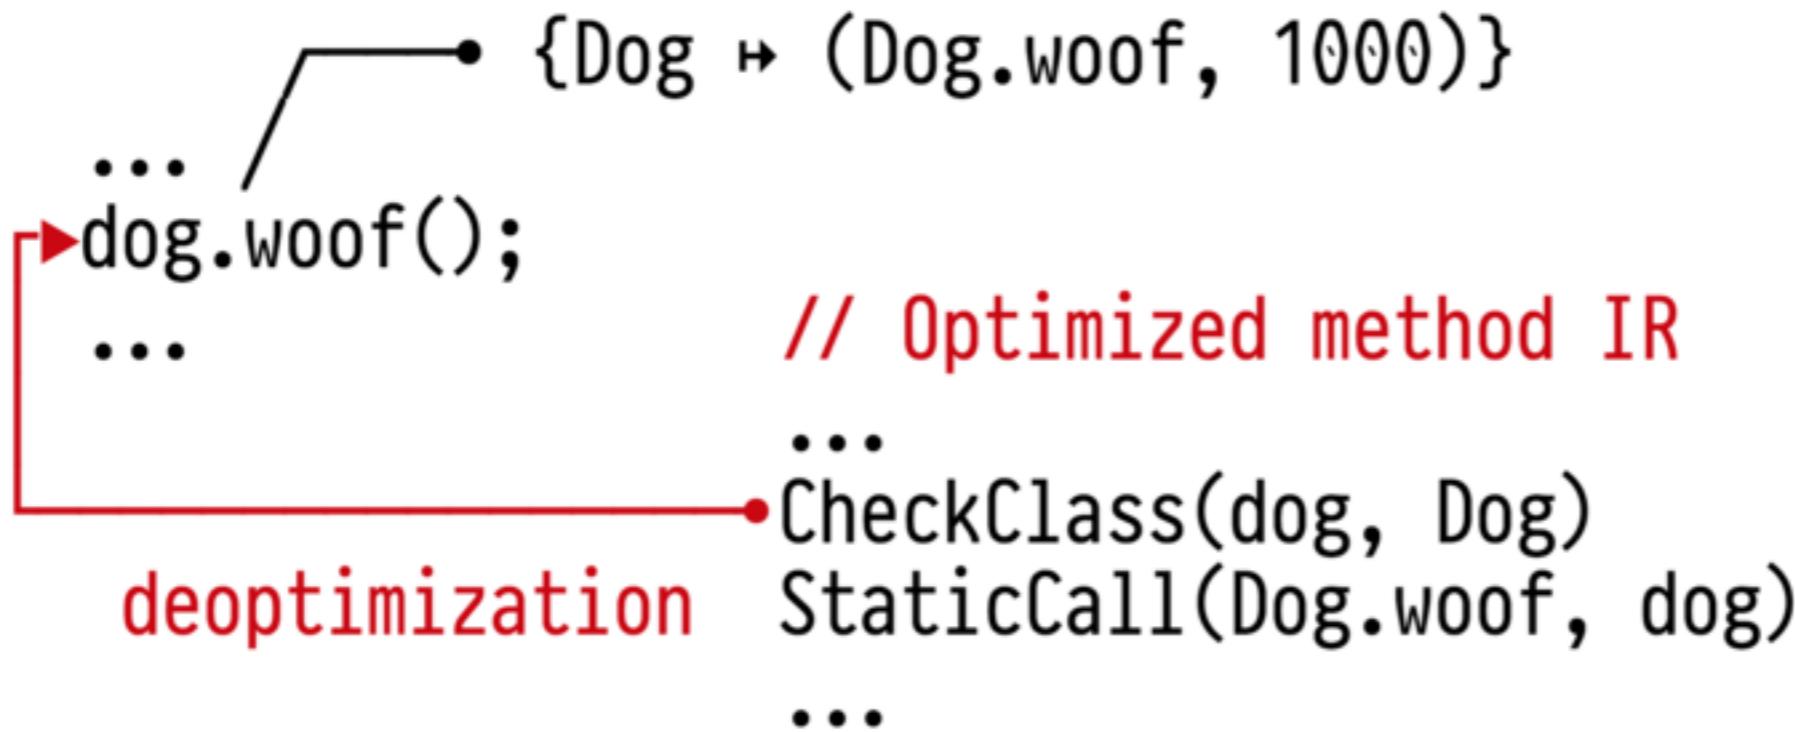
```
    └─ {Dog ↪ (Dog.woof, 4)}  
        ↴ LookupAndCache() ↑  
dog.woof();  
{  
    └─ class Wolf {  
        woof() => print("🐺🐺🐺");  
    }  
}
```

```
    ↗ {Dog ↪ (Dog.woof, 4),  
     Wolf ↪ (Wolf.woof, 1)}  
dog.woof();  
{  
  ↘ class Wolf {  
    woof() => print("🐺🐺🐺");  
  }  
}
```

type profiling

```
... └─ {Dog ↪ (Dog.woof, 1000)}  
... dog.woof();  
...
```

```
... {Dog ↪ (Dog.woof, 1000)}  
... dog.woof(); ...  
...  
// Optimized method IR  
...  
CheckClass(dog, Dog)  
StaticCall(Dog.woof, dog)  
...
```



classics

Optimizing Dynamically-Typed Object-Oriented Languages With Polymorphic Inline Caches

Urs Hözle
Craig Chambers
David Ungar[†]

Computer Systems Laboratory, Stanford University, Stanford, CA 94305
{urs,craig,ungar}@self.stanford.edu

Abstract: *Polymorphic inline caches* (PICs) provide a new way to reduce the overhead of polymorphic message sends by extending inline caches to include more than one cached lookup result per call site. For a set of typical object-oriented SELF programs, PICs achieve a median speedup of 11%.

As an important side effect, PICs collect type information by recording all of the receiver types actually used at a given call site. The compiler can exploit this type information to generate better code when *recompiling* a method. An experimental version of such a system achieves a median speedup of 27% for our set of SELF programs, reducing the number of non-inlined message sends by a factor of two.

Implementations of dynamically-typed object-oriented languages have been limited by the paucity of type information available to the compiler. The abundance of the type information provided by PICs suggests a new compilation approach for these languages, *adaptive compilation*. Such compilers may succeed in generating very efficient code for the time-critical parts of a program without incurring distracting compilation pauses.

1. Introduction

Historically, dynamically-typed object-oriented languages have run much slower than statically-typed languages. This disparity in performance stemmed largely from the relatively slow speed and high



surprisingly
general

```
dog.woof(); // Looks like a call, but is actually
{           // a field load and a call.

    class Alien {
        var woof; // A field
        Alien() : woof = () => print("👽👽👽");
    }
}
```

```
//      (...) => dog.woof(...)  
var f = dog.woof;  
  
method tear-off
```

```
dog.woof();  
{  
    class Cat {  
        noSuchMethod(invocation) {  
            if (invocation.memberName == #woof) {  
                print('I am a 🐱 not a 🐶!');  
            }  
        }  
    }  
}
```

```
{}  
dog.woof();  
  
class Cat {  
    noSuchMethod(invocation) {  
        if (invocation.memberName == #woof) {  
            print('I am a 🐱 not a 🐶!');  
        }  
    }  
}
```

```
{}  
    ↴ LookupAndCache()  
dog.woof();  
  
class Cat {  
    noSuchMethod(invocation) {  
        if (invocation.memberName == #woof) {  
            print('I am a 🐱 not a 🐶!');  
        }  
    }  
}  
}
```

```
{Cat ↪ ( ?, 1)}  
↳ LookupAndCache()  
dog.woof();  
  
class Cat {  
    noSuchMethod(invocation) {  
        if (invocation.memberName == #woof) {  
            print('I am a 🐱 not a 🐶!');  
        }  
    }  
}  
}
```

```
{}  
dog.woof();  
  
class Cat {  
    noSuchMethod(invocation) {  
        if (invocation.memberName == #woof) {  
            print('I am a 🐱 not a 🐶!');  
        }  
    }  
}
```

```
{}  
    ↴ LookupAndCache()  
dog.woof();  
  
class Cat {  
    noSuchMethod(invocation) {  
        if (invocation.memberName == #woof) {  
            print('I am a 🐱 not a 🐶!');  
        }  
    }  
}  
}
```

```
{}  
    ↴ LookupAndCache()  
dog.woof();  
{  
class Cat {  
    noSuchMethod(invocation) {  
        if (invocation.memberName == #woof) {  
            print('I am a 🐱 not a 🐶!');  
        }  
    }  
}  
}
```

```
{}  
dog.woof();  
  
class Cat {  
    noSuchMethod(invocation) {  
        if (invocation.memberName == #woof) {  
            print('I am a 🐱 not a 🐶!');  
        }  
    }  
}
```

```
{}  
    ↪ LookupAndCache()  
dog.woof();  
{  
    class Cat {  
}  
}
```

```
{}  
    ↳ LookupAndCache()  
dog.woof();  
{  
    class Cat {  
        // ↓ injected  
        noSuchMethod:woof() {  
            return noSuchMethod(  
                Invocation(memberName: #woof));  
        }  
    }  
}
```

```
{}  
↳ LookupAndCache()  
dog.woof();  
  
class Cat {  
    // ↓ injected  
    noSuchMethod:woof() {  
        return noSuchMethod(  
            Invocation(memberName: #woof));  
    }  
}
```

```
{ Cat ↪ (Cat.noSuchMethod:woof, 1) }
```

dog.woof();

```
{
```

```
class Cat {  
    // ↓ injected  
    noSuchMethod:woof() {  
        return noSuchMethod(  
            Invocation(memberName: #woof));  
    }  
}
```

suddenly
all optimizations
just work

2012 - 2013

new JIT is "done"

non-classical

```
... ┌─────────┐ {Dog ↪ (Dog.woof, 1000)}  
... dog.woof(); └─────────┘  
...  
// Optimized method IR  
...  
CheckClass(dog, Dog)  
StaticCall(Dog.woof, dog)  
...
```

```
... {Dog ↪ (Dog.woof, 1000)}  
... dog.woof(); ...  
...  
// Optimized method IR  
...  
CheckClass(dog, Dog) ← check at use  
StaticCall(Dog.woof, dog)  
...
```

Globally Tracked Field State

```
class DogHolder {  
    var dog = Dog();  
    final listOfDogs = List<Dog>(10);  
}
```

The diagram illustrates the state tracking for fields in the `DogHolder` class. It uses red arrows to point from specific field declarations to annotations. The first arrow points from the declaration of `dog` to the annotation "(class, nullability, length-of-array)". The second arrow points from the declaration of `listOfDogs` to the same annotation.

- Loads can *assume* the state
[Concrete classes are used to eliminate CheckClass, length is is used to eliminate bounds checks.]
- Stores update the state
[If incompatible change: optimized code depending on the assumption is deoptimized]

Globally Tracked Field State

```
class DogHolder {  
    var dog = Dog();  
  
    final listOfDogs = List<Dog>(10);  
}
```

The code snippet shows three tracked fields:

- `dog`: An annotation with a red line pointing to it: **(Dog, not-null, N/A)**.
- `listOfDogs`: An annotation with a red line pointing to it: **(List, not-null, 10)**.

Optimization shifts checks from uses to definitions

[where they often can be eliminated]

Strange Loop 2013

« Building an Optimizing Compiler for Dart »

2014

tweaking things

« to hand-written assembly» effort

intrinsics rewritten in IL

[planned to move stubs too, still not started on that]

irregexp port

- irregexp is a V8's JITing RE engine
- in V8 it has platform specific backends
- in Dart we just generate IL from RegExp-s

• • •

« Wait, I thought Dart is
supposed to be a Web
language 🤔 »

VM

dart2js

VM

dart2js

Parallel evolution of Dart-to-JS compilers

- **dartc** [AST based written in Java]
- **frog** [AST based "transpiler" in Dart]
- **leg** [SSA based written in Dart]
- **dart2js** [Actually just renamed leg]

AOT compilation
of Dart to JS
is hard

Object model mismatch

can not completely erase Dart object model in JS
representation only as efficient as JS engine allows it



Open the Sky

[Browse files](#)

git master ⌂ fitness_apk_4 ... demo_apk_22

 abarth committed on Oct 23, 2014

1 parent 20cc569 commit ae72930937b1a1522d1919a82ddc21da7d953229

 Showing 4,395 changed files with 688,796 additions and 0 deletions.

[Unified](#) [Split](#)

The diff you're trying to view is too large. We only load the first 3000 changed files.

«What would happen if you
take WebView and remove the
constraint that it has to
render web-pages?»

— Adam Barth at FLOSS Weekly 439

Initially Sky started with JS
but later moved to Dart *entirely*

[with C++ for lowest *engine* layers]

and renamed



talk from Strange Loop 2017

« **Flutter: How we're building
a UI framework for tomorrow
at Google» by Eric Seidel**



Flutter
is a mobile framework

iOS does not allow JITing

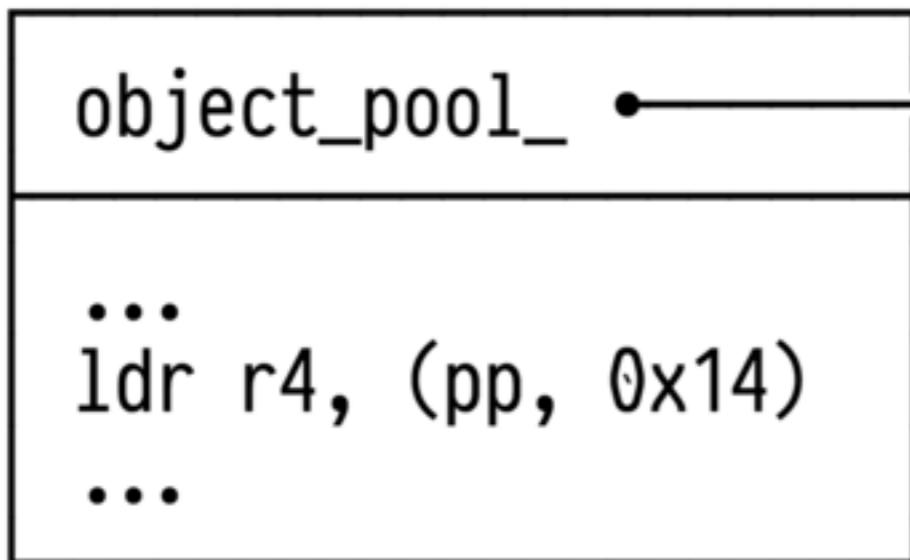
[... unless you are Mobile Safari]

AOT

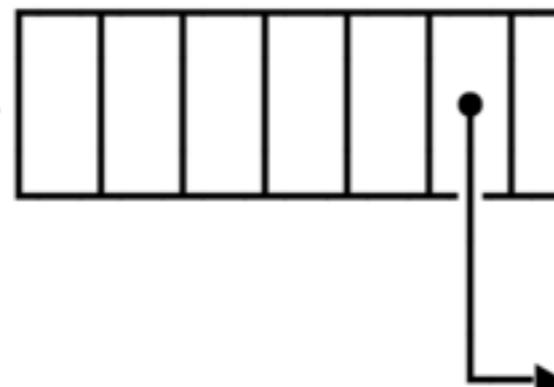
2015

started AOT prototype

instructions



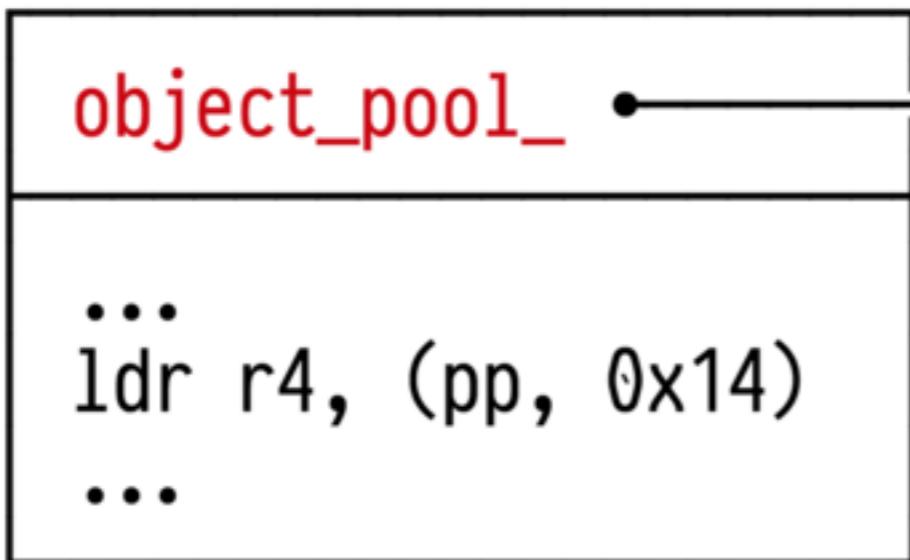
pool



object
in the heap

JIT code
is relocatable

instructions



pool

object
in the heap

had to change calling conventions

[want instructions to be RO and *isolate* independent, while pools are isolate dependent - so instructions can't contain pointer to the pool]

« don't want flags & no-body cares about JIT »

[changed calling conventions in a way that impacted JIT performance — in retrospect was a mistake]

then just
serialize it!

already had
snapshots

[kinda like a heap dump that you can hydrate back]

how to get "JIT" code?

[real JIT does speculative opts!]

disable speculations then
force compile everything

DONE

[prototype took around 2-3 months]

"DONE"

[...]

"DONE"

[...]

AOT compilation
of Dart to JS
is hard

AOT compilation
of Dart to []

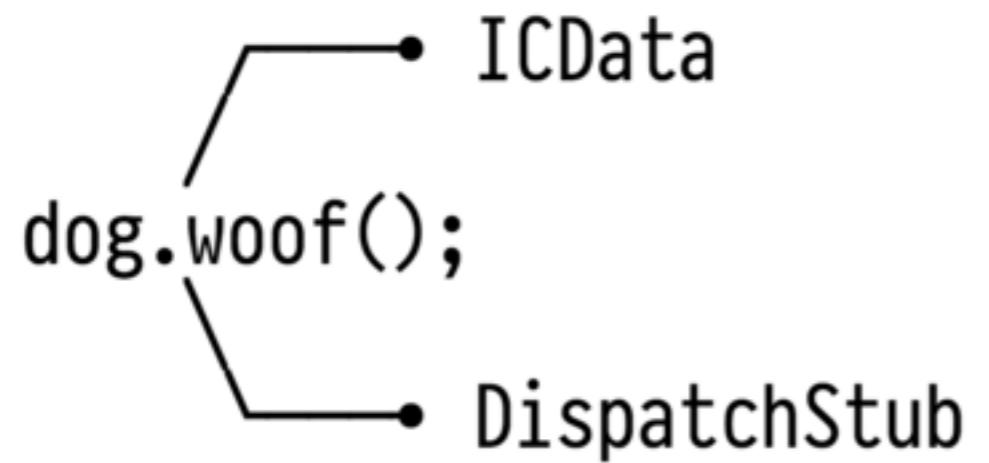
is nontrivial

need replacement
for **SPECULATIONS**

switchable calls

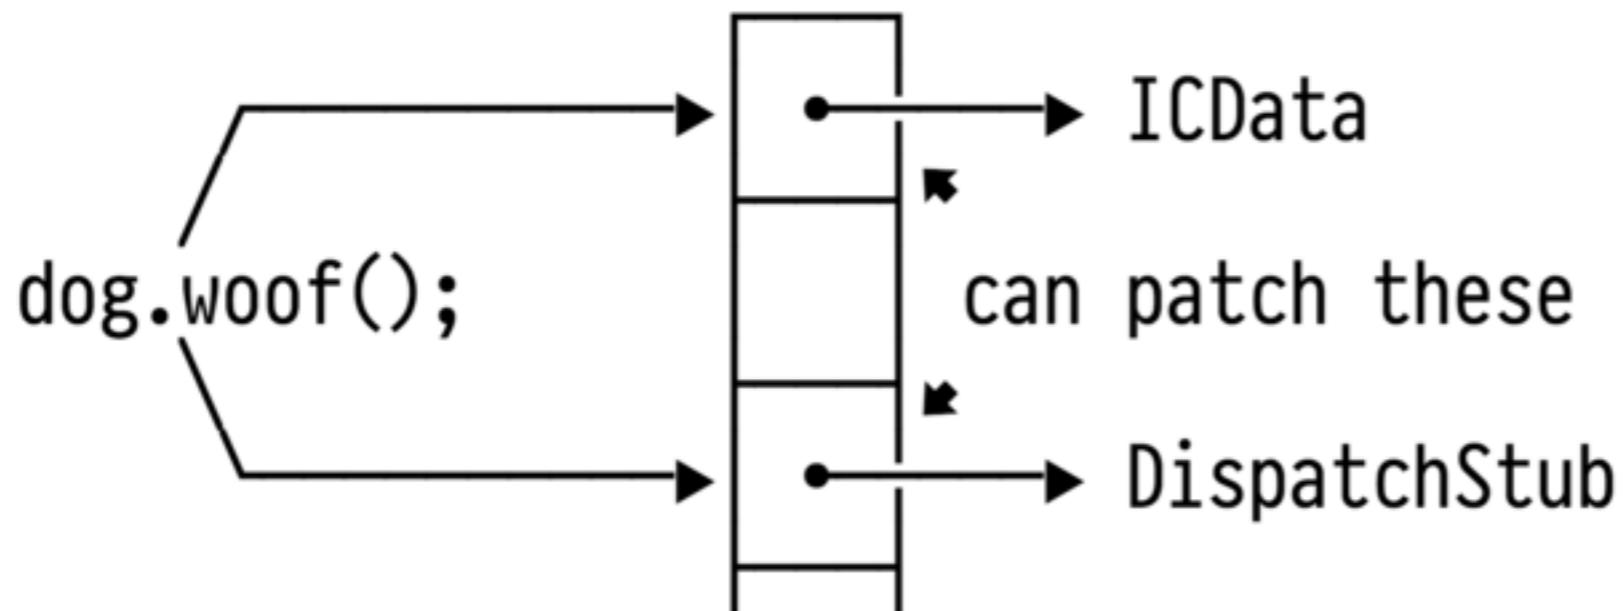
```
// call-site specific cache
ICData { ... }

dog.woof();
// generic dispatch stub
dispatch(icData, receiver, ...) { ... }
```



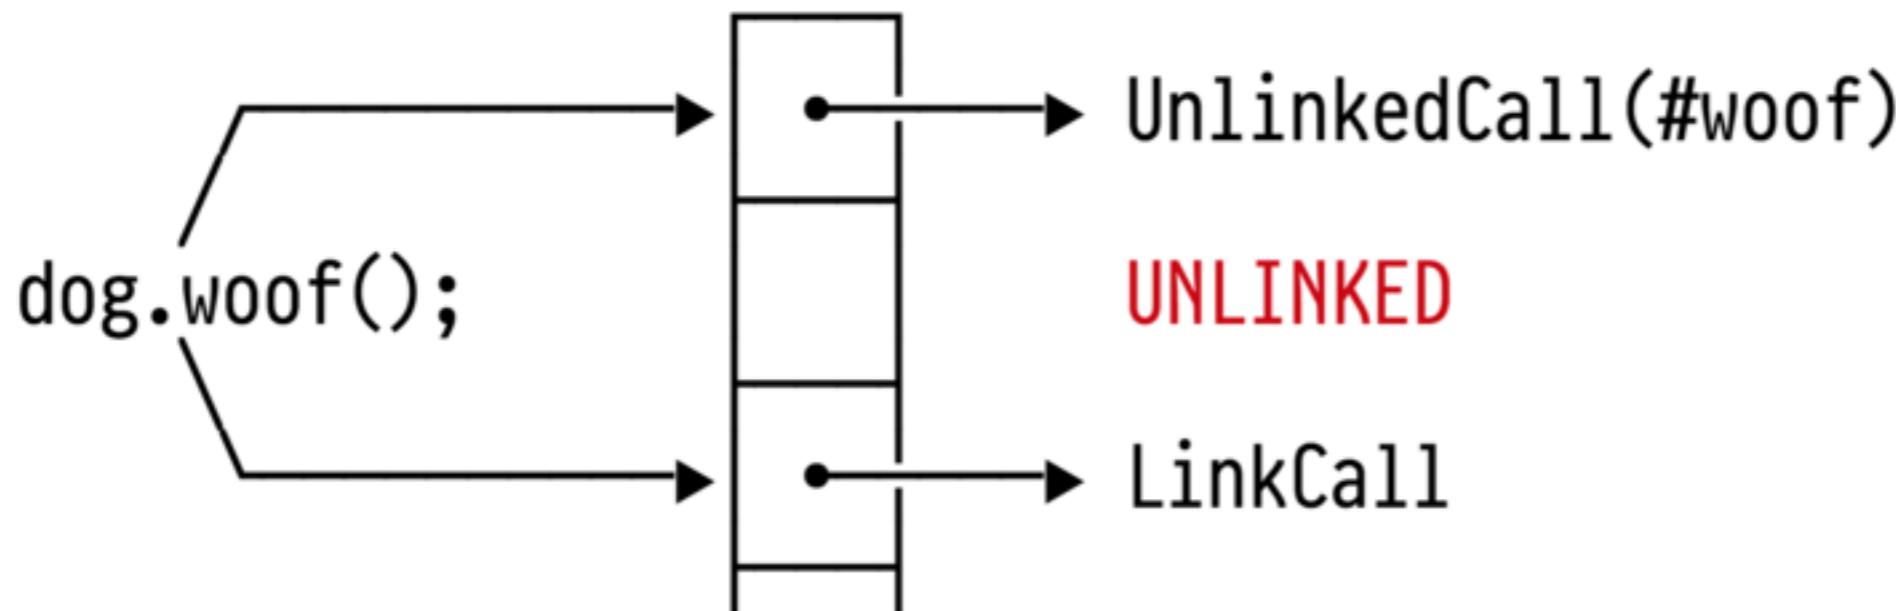
code (rx)

pool (rw)



code (rx)

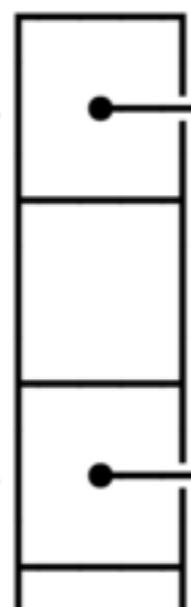
pool (rw)



code (rx)

dog.woof();

pool (rw)



class Dog

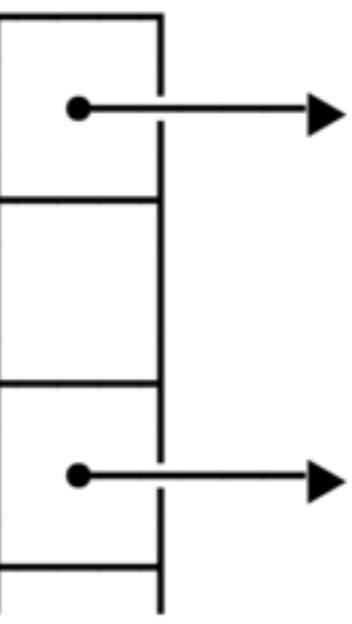
MONOMORPHIC

Dog.woof#monomorphicEntry

code (rx)

dog.woof();

pool (rw)



class Dog

Dog.woof

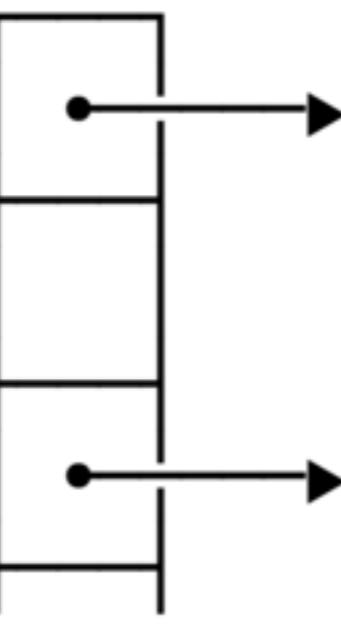
normal
entry

```
if (this.class != expected)  
    return monomorphicMiss(...);
```

code (rx)

dog.woof();

pool (rw)



class Dog

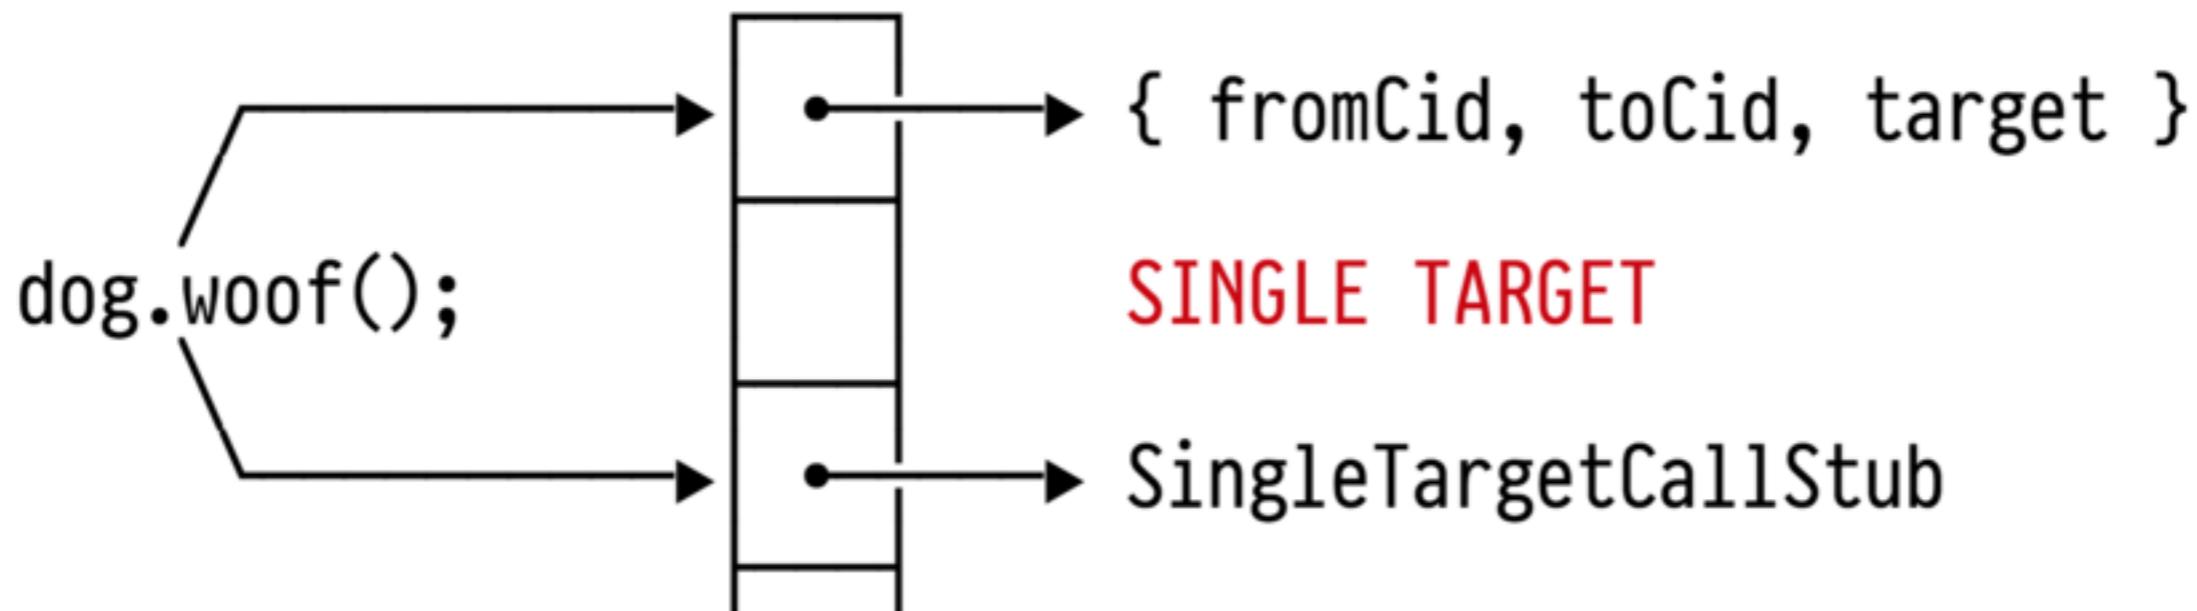
Dog.woof

normal
entry

```
if (this.class != icData)  
    return monomorphicMiss(...);
```

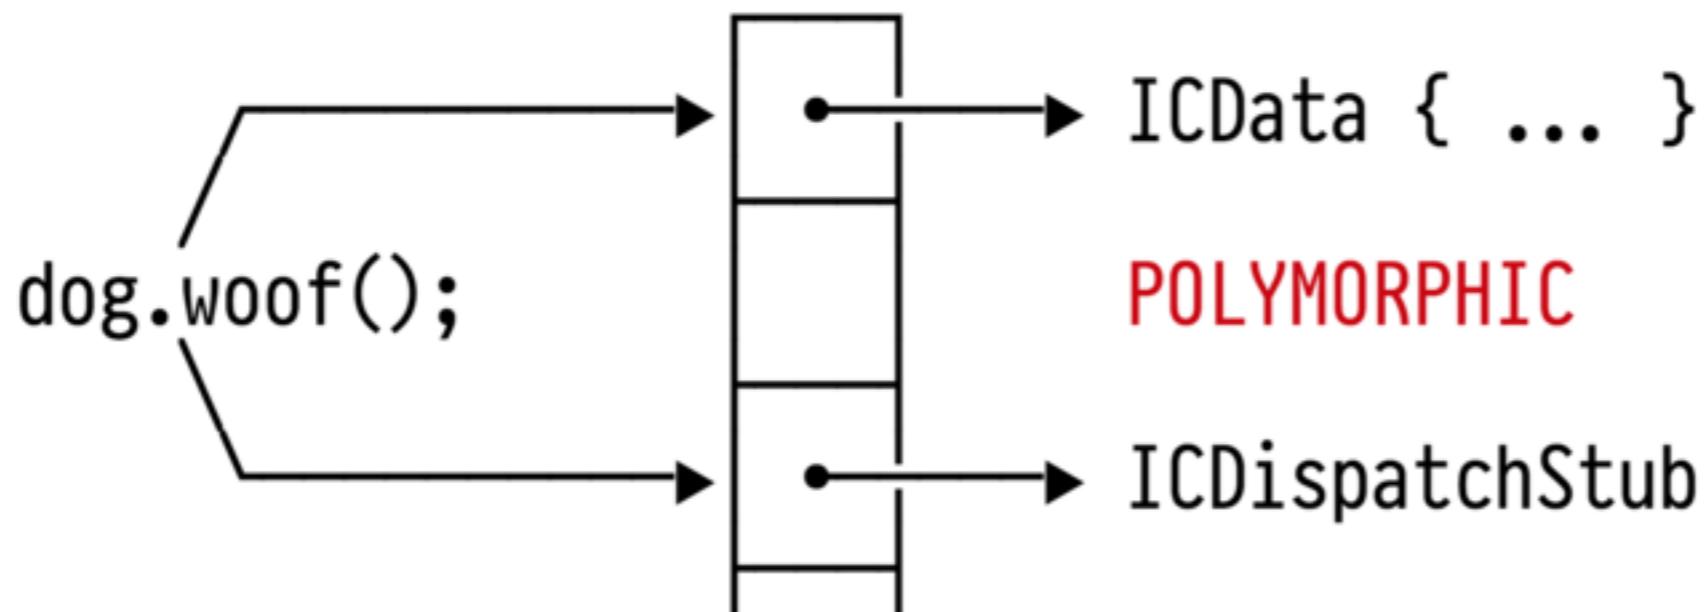
code (rx)

pool (rw)



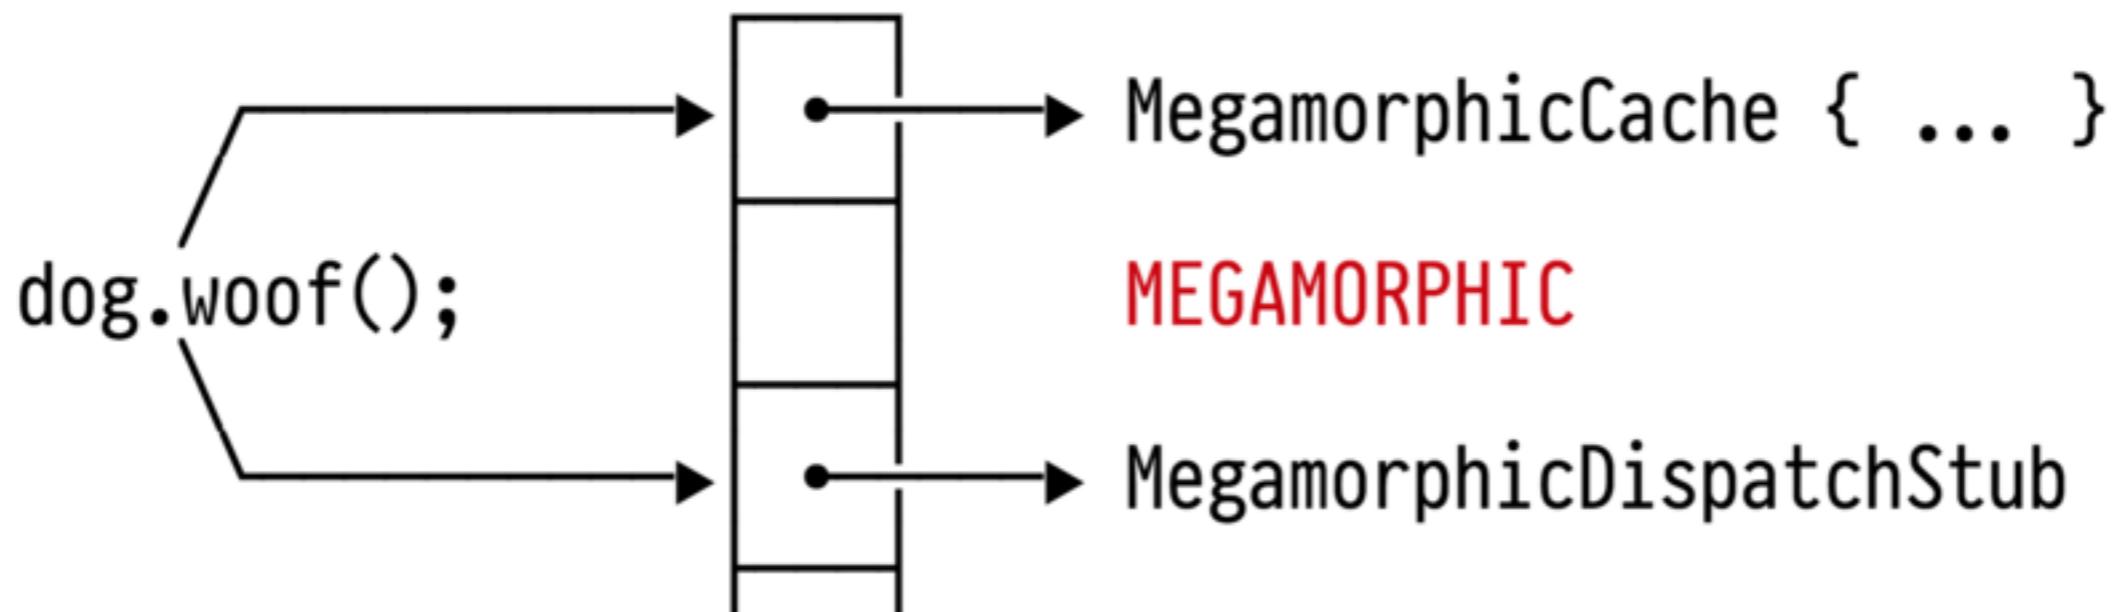
code (rx)

pool (rw)



code (rx)

pool (rw)



inline arithmetic fast-paths

- operators can be overridden
- int is nullable in Dart
- int was arbitrary precision

```
; ; Y ← CheckedSmiOp(+, X, 1)
testl X, 1
jnz slow•
movq Y, X
addq Y, 2
jo slow•
► done:
...
slow: ;; "cold" section of the code
pushq X
pushq 1
movq RBX, [PP+...] ;; megamorphic cache
callq THR→megamorphic_call_stub
movq Y, RAX
jmp done
```

```
; ; Y ← CheckedSmiOp(+, X, 1)
testl X, 1    ;; what is this if (X & 1 ≠ 0) check?
jnz slow•
movq Y, X
addq Y, 2    ;; why are we adding 2?
jo slow•
```

► done:

...

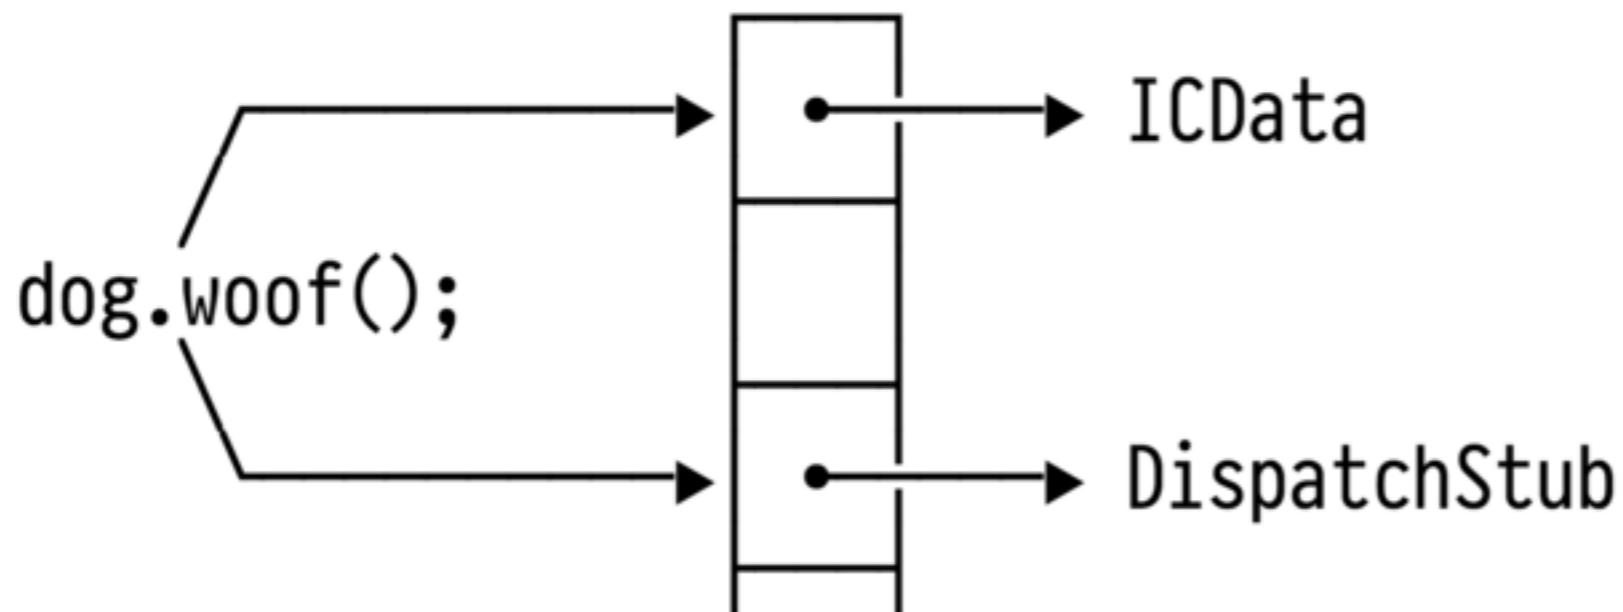
```
slow:  ;; "cold" section of the code
pushq X
pushq 1
movq RBX, [PP+...] ;; megamorphic cache
callq THR→megamorphic_call_stub
movq Y, RAX
jmp done
```

plus some
closed-world
CHA opts

what about
code size?

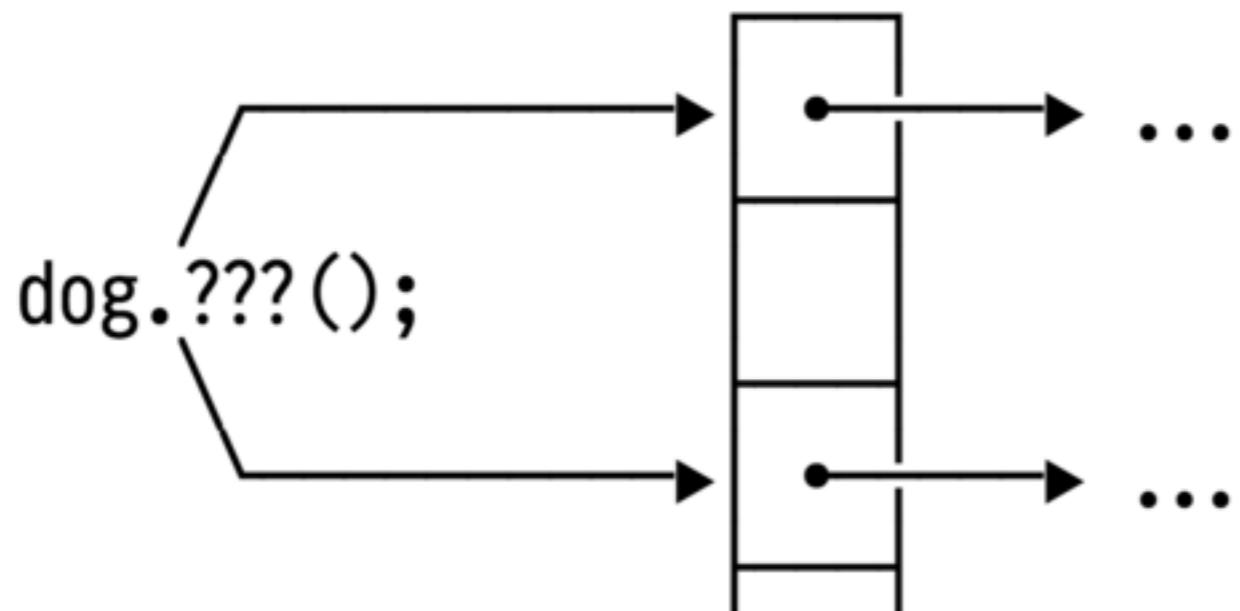
code (rx)

pool (rw)



code (rx)

pool (rw)



```
void doOne(x) {  
    return x.foo().bar();  
}
```

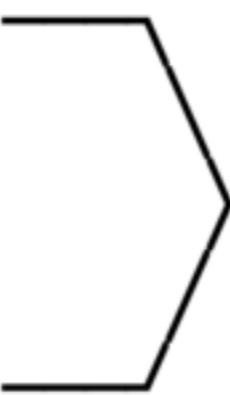
```
void doTwo(z) {  
    return z.baz().quux();  
}
```

```
void doOne(p0) {  
    return p0.???.().???.();  
}
```

```
void doTwo(p0) {  
    return p0.???.().???.();  
}
```

```
void doOne(p0) {  
    return p0.???.().??();  
}
```

```
void doTwo(p0) {  
    return p0.???.().??();  
}
```



same machine code
and can be shared

byproduct of AOT work

AppJIT snapshots

[warmup the application - then serialize JIT code; massive reduction in warmup time — used by our tools now]

only so much can be done

[remember for example dynamic method injection?]

2016

the wind of changes

(1) dart2js is too slow
for development

Dart Dev Compiler

has been cooking since 2014

Dart Dev Compiler

fast-modular compiles

Dart Dev Compiler

requires stronger type system

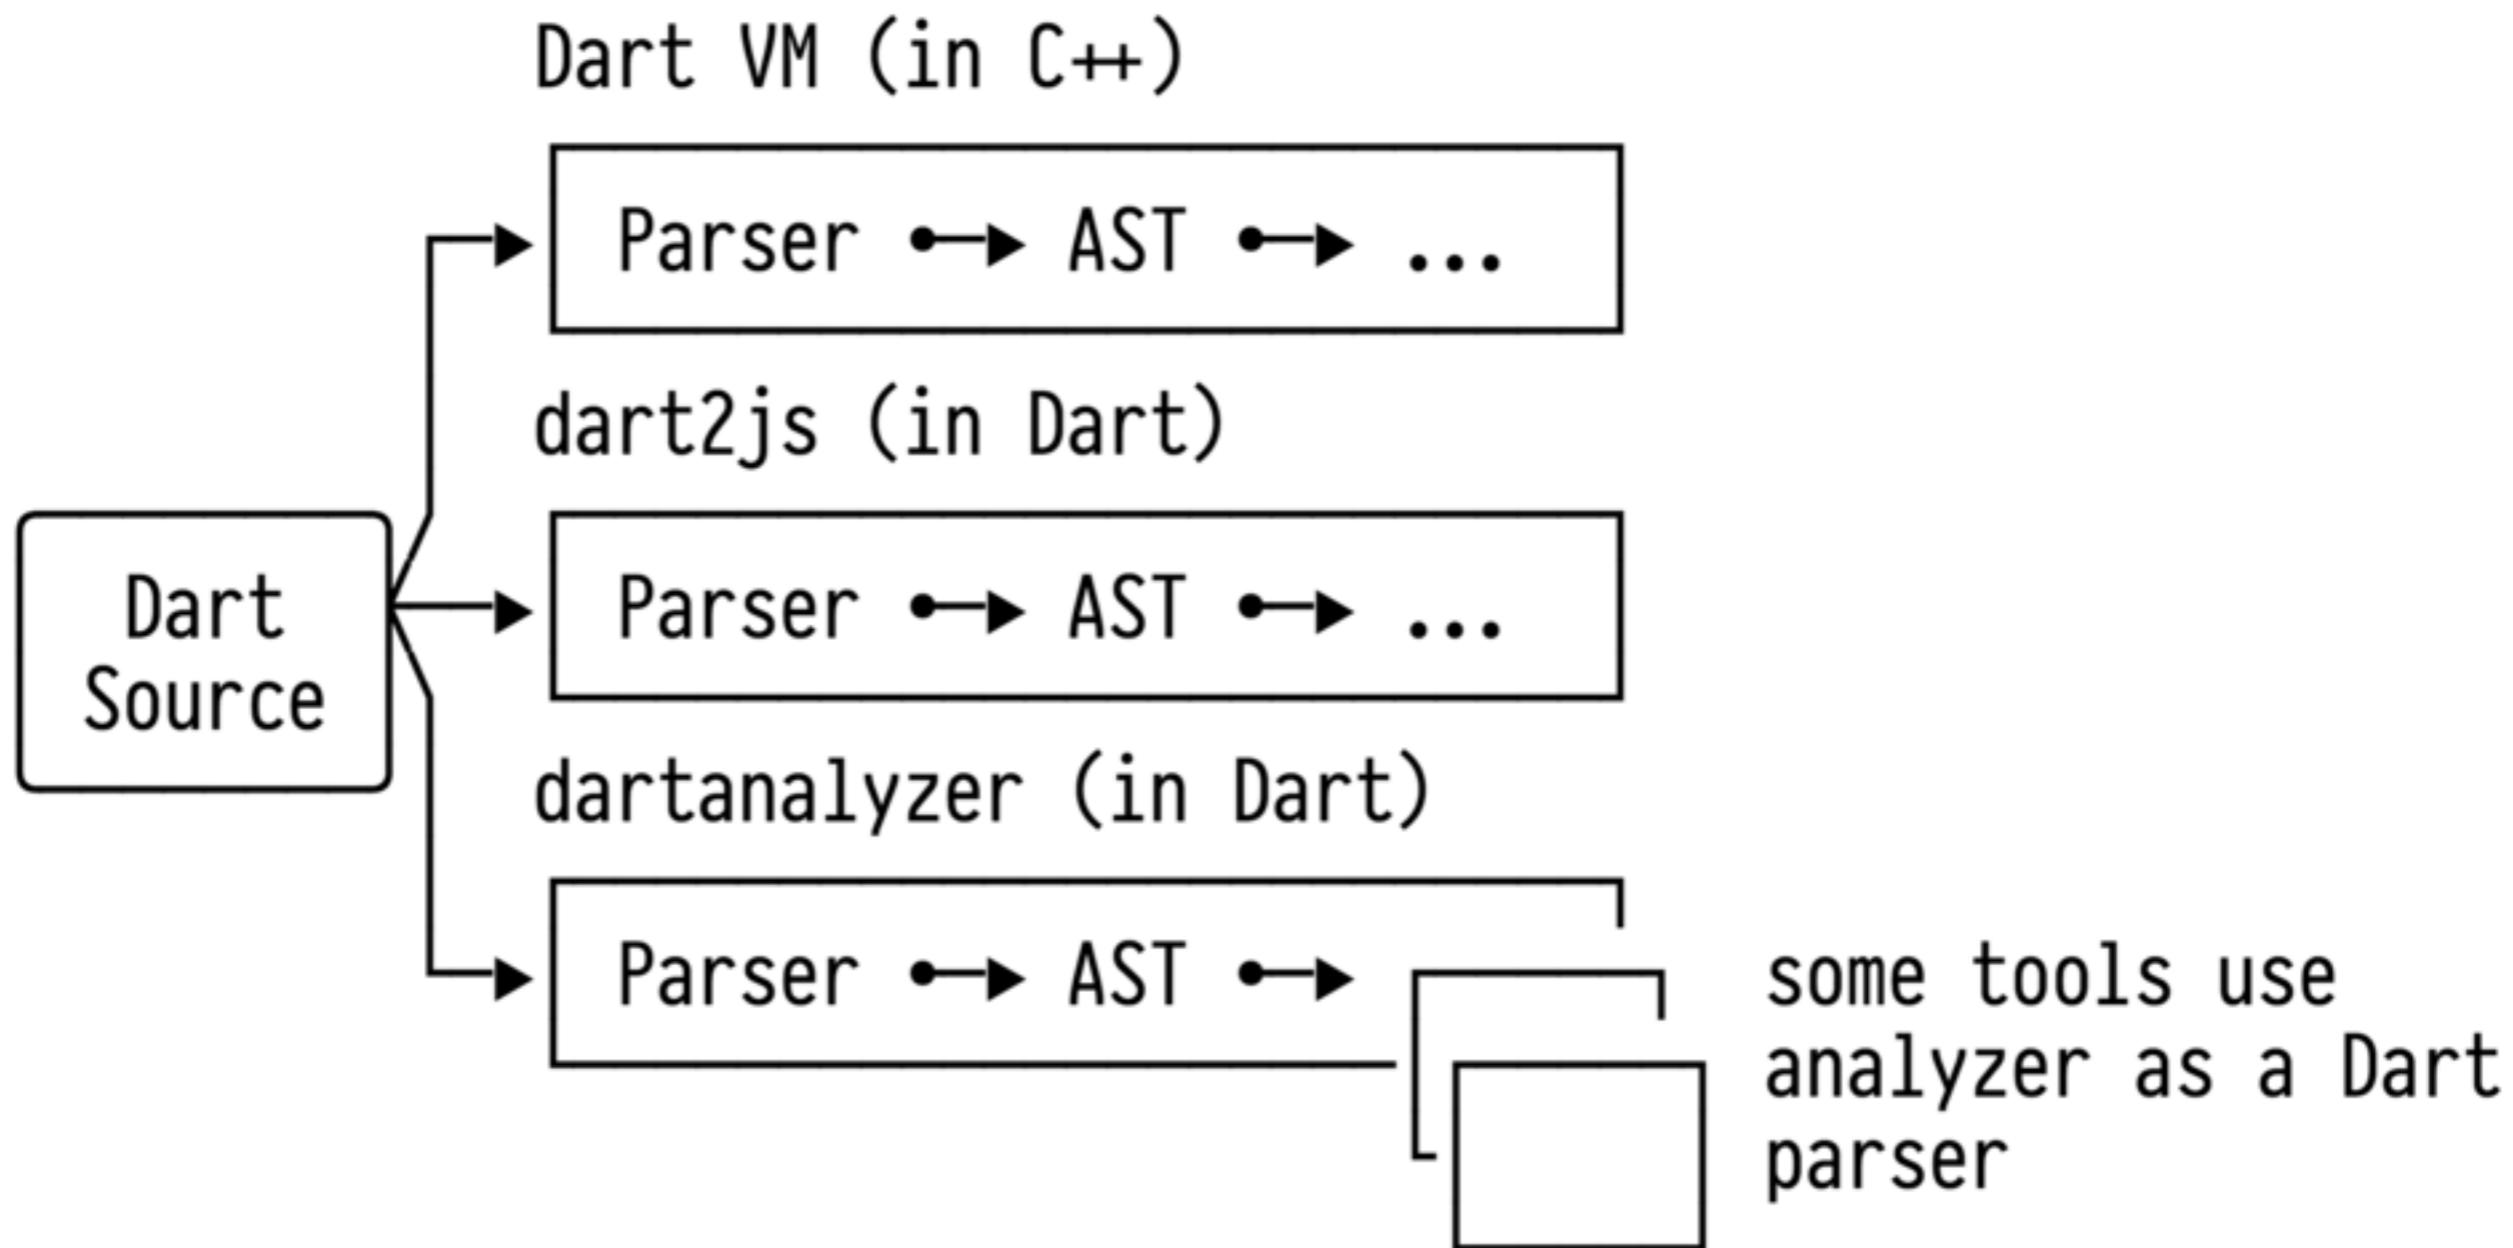
```
Dog d; // is guaranteed to be a 🐶
```

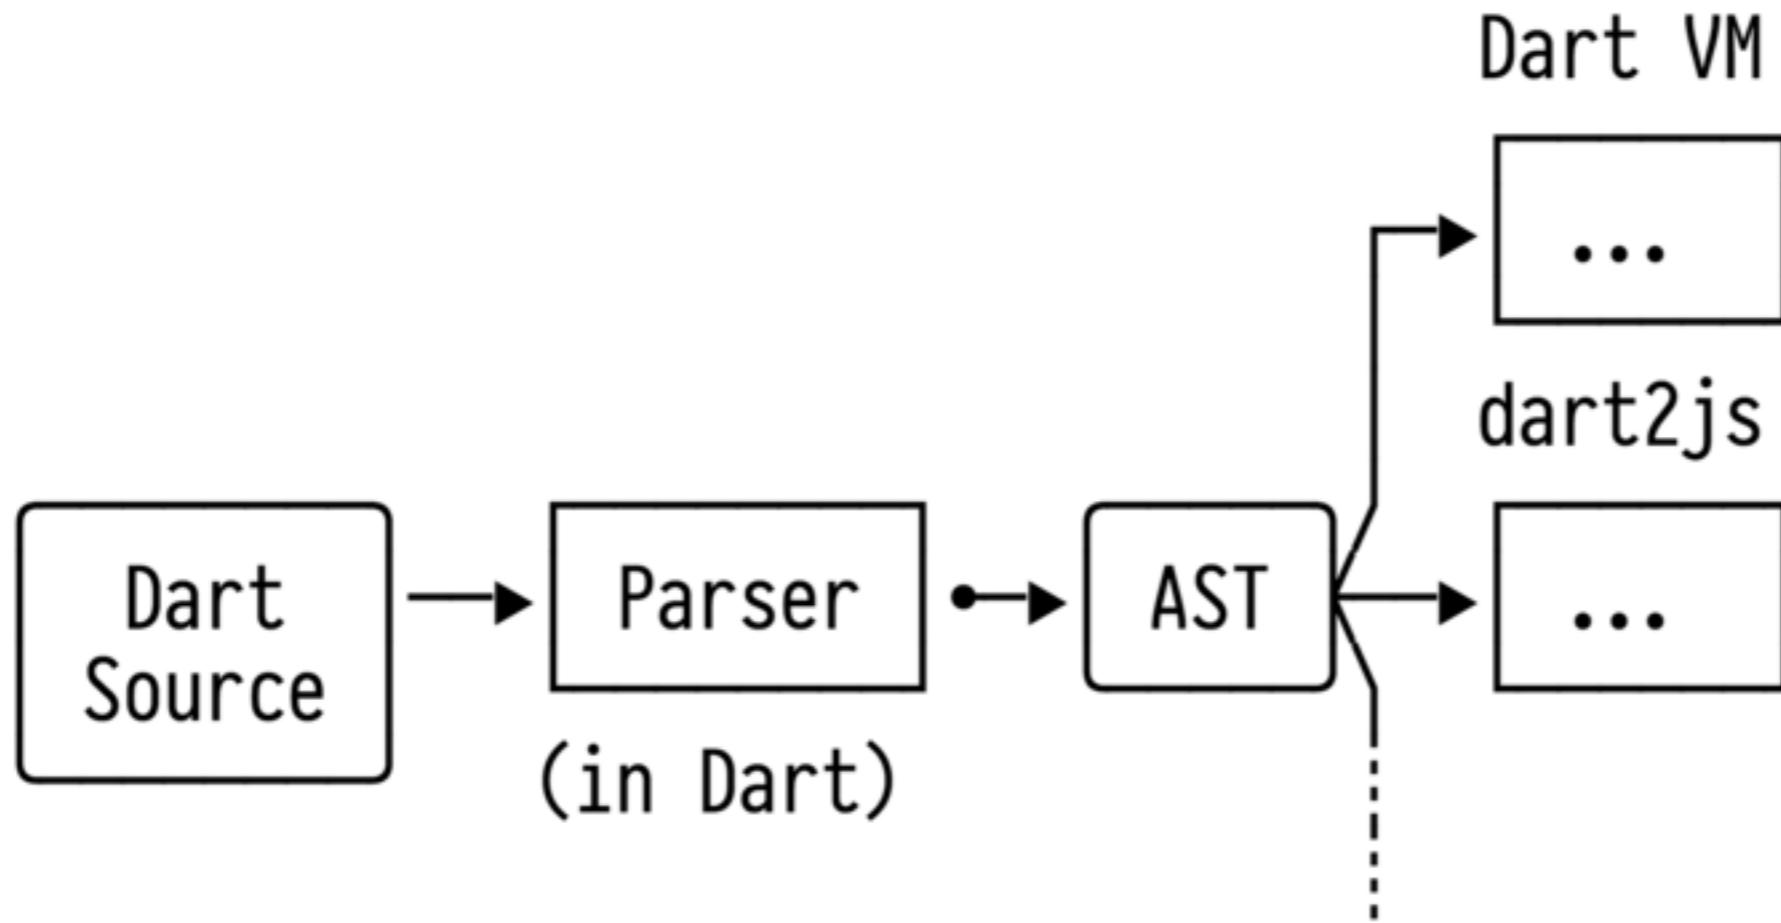
```
List<Cat> x;  
x[0]; // is guaranteed to be a 🐱
```

talk from Dart Conf 2018

« Evolving Dart: Leaving the
ocean and learning to fly »
by Leaf Petersen

(2) every Dart tool
has its own parser





2016

new type system
& common front-end
migration starts

Aug 2018

Dart 2 shipped

tons of *gnarly* work

[cool stuff too: e.g. started implementing optimization passes in Dart]

is AOT trivial now?

- many important types are *interfaces*, e.g. `List<T>`
- `int` and `double` are still nullable
- `int` became 64-bit wrap around integer
- *soundness* requires runtime type checks

```
// In Dart all generics are covariant
var strings = ["a", "b"]; // static type List<String>
strings.add(10); // compile time error
```

```
// Upcast List<String> to List<Object>
List<Object> objects = strings; // ok
objects.add(10); // runtime error
// requires check somewhere
```

```
// In Dart all generics are covariant
var strings = ["a", "b"]; // static type List<String>
strings.add(10); // compile time error
```

```
// Upcast List<String> to List<Object>
List<Object> objects = strings; // ok
objects.add(10); // runtime error
// requires check somewhere
```

```
// Also Dart 2 still has dynamic
dynamic x = strings;
x.append("a"); // runtime error
x.add(10); // runtime error
```

Current approach:

- track type arguments (statically in AOT, dynamically in JIT)
- bypass type checks when runtime type arguments are known to match static type arguments

List<T>.join(String sep)

arguments not affected by covariance – no checks

► every call enters here

List<T>.join(String sep)

arguments not affected by covariance – no checks

► every call enters here
except dynamic ones

```
dynamic x = [...];  
x.join(10);
```

List<T>.join(String sep)

arguments not affected by
covariance – no checks

```
dynamic x = [...];  
x.dyn:join(10);
```

List<T>.join(String sep)

arguments not affected by
covariance – no checks

List<T>.dyn:join(String sep)

Assert(sep is String)
Call List.join •

```
dynamic x = [...];  
x.join(10);
```

List<T>.join(String sep)

arguments not affected by
covariance – no checks

List<T>.add(T elem)

arguments are affected by covariance – need to check

List<T>.add(T elem)

Assert(elem is T)

- checked entry
- unchecked entry

```
    } match {  
        List<String> strs = new List<String>();  
        strs.add(x);  
    }
```

```
List<T>.add(T elem)
```

```
    Assert(elem is T)
```

↳ checked entry

↳ unchecked entry

not match

```
List<Object> strs = new List<String>();  
strs.add(x);
```

List<T>.add(T elem)

Assert(elem is T)

- checked entry
- unchecked entry

can use static analysis or dynamic profiling
to determine / speculate invariance

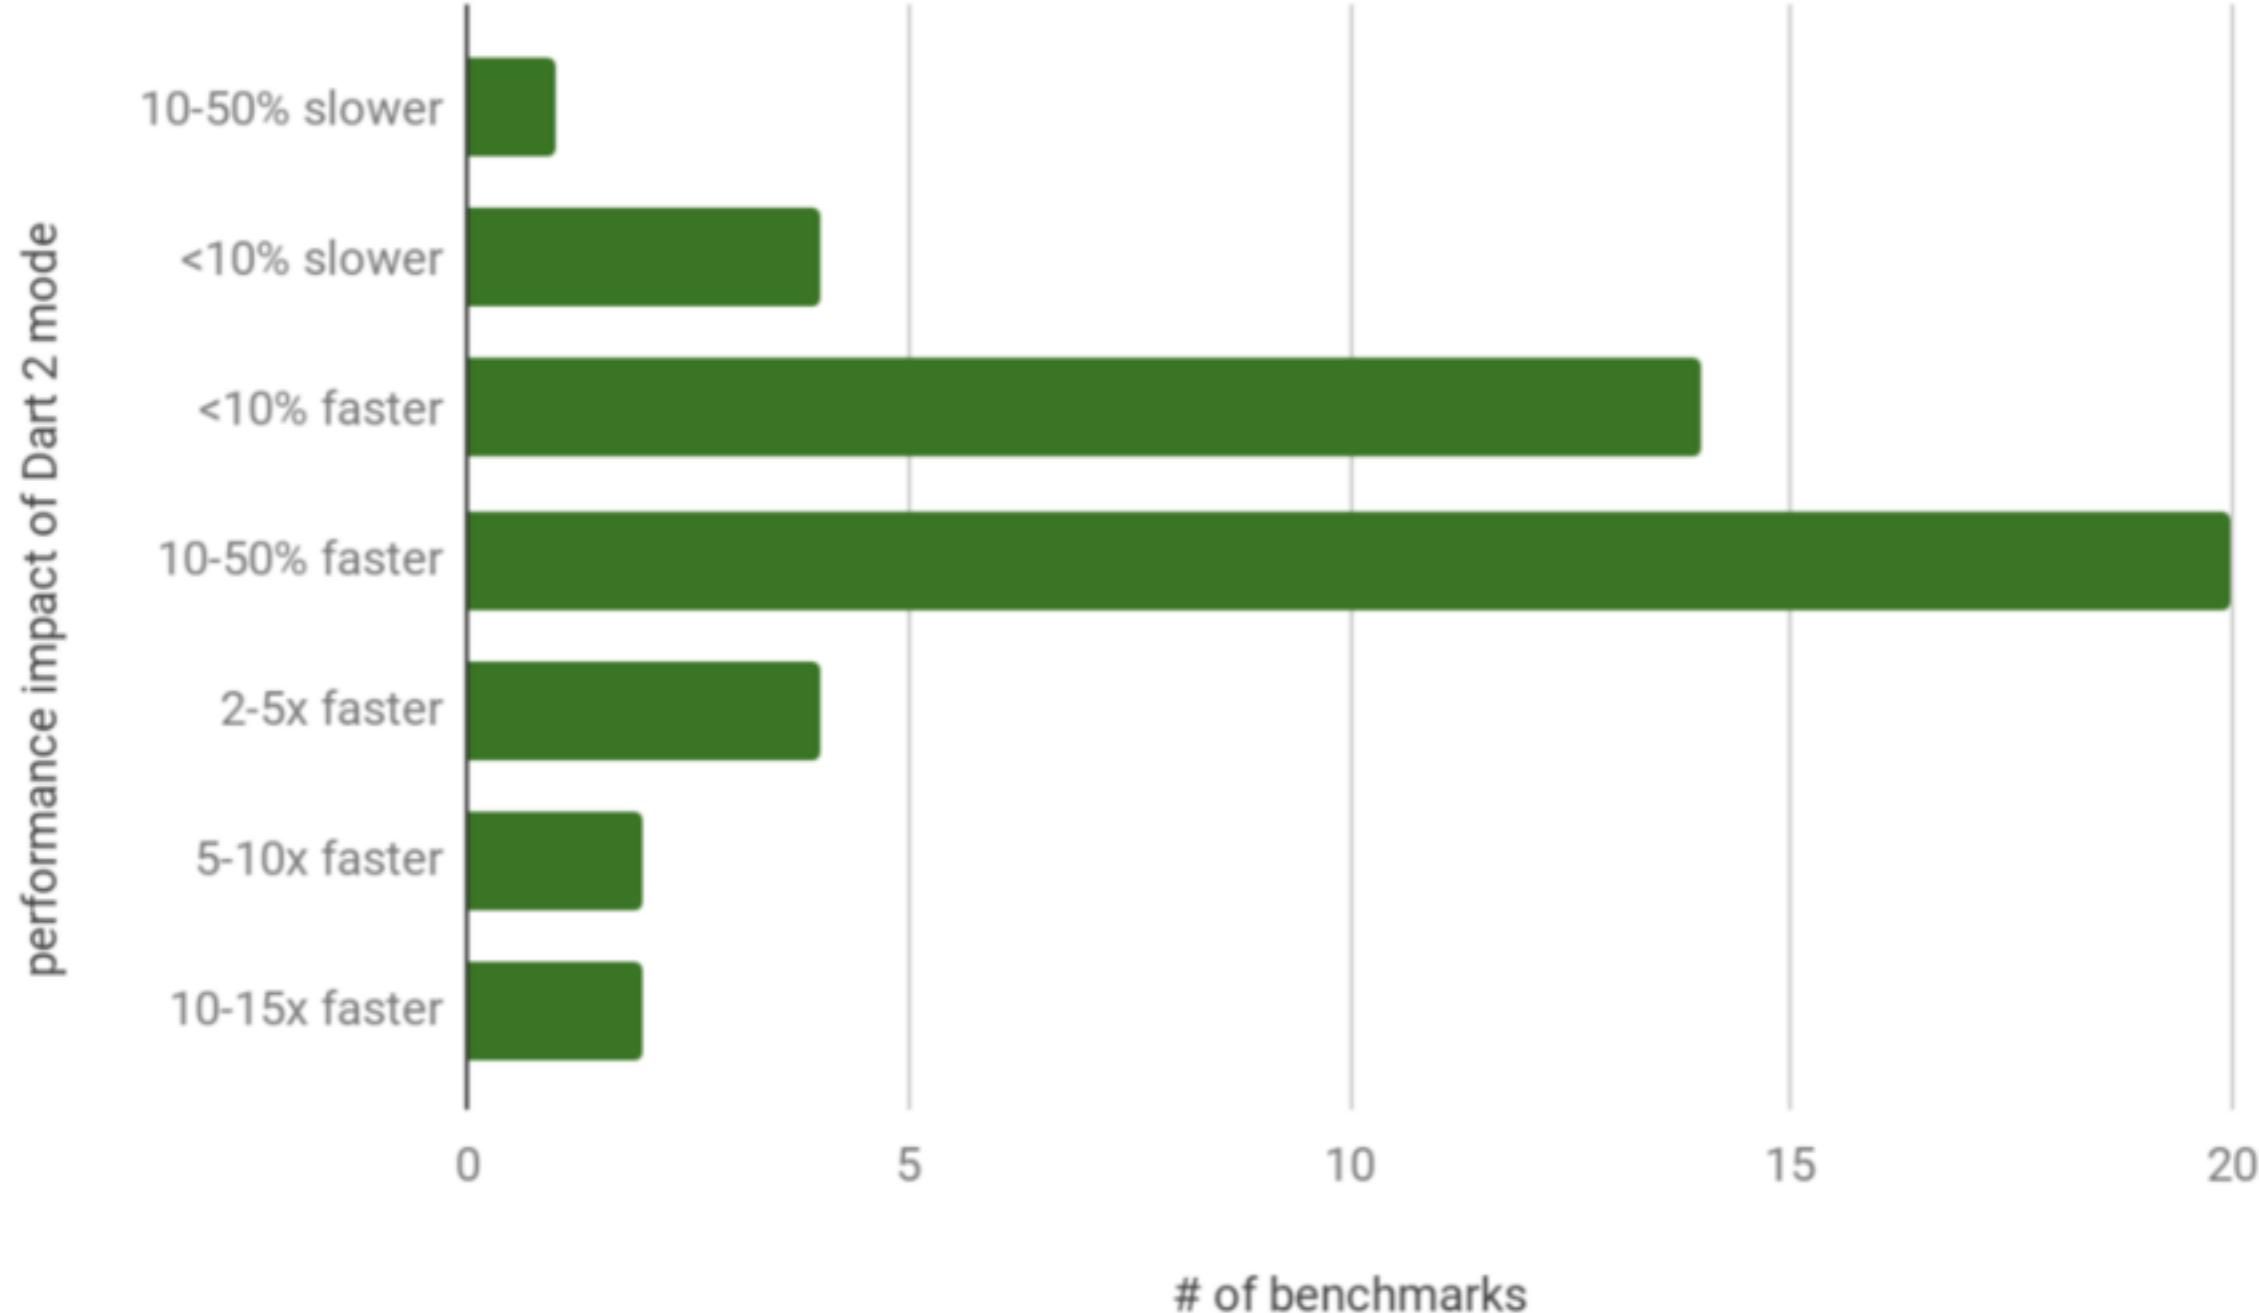
```
List<String> strs = ...;  
strs.add(x);
```

```
List<T>.add(T elem)
```

```
Assert(elem is T)
```

► checked entry

► unchecked entry



architectural changes
take forever
[choose the balance wisely]

parameterize your
architecture

optimizations are more
general than you think

[just need to tweak things a bit]

there are still interesting
problems to solve

[code size for example]

if you have your own
compiler use it more
[don't write assembly by hand]